

sprawozdanie

June 4, 2023

1 GPU -CUDA

Wykonał: Paweł Kruczkiewicz

Wykonano w środowisku Google Collab na Tesli T4

1.1 Ćw 1

Kod:

```
/* Copyright (c) 2022, NVIDIA CORPORATION. All rights reserved.
 *
 * Redistribution and use in source and binary forms, with or without
 * modification, are permitted provided that the following conditions
 * are met:
 *   * Redistributions of source code must retain the above copyright
 *     notice, this list of conditions and the following disclaimer.
 *   * Redistributions in binary form must reproduce the above copyright
 *     notice, this list of conditions and the following disclaimer in the
 *     documentation and/or other materials provided with the distribution.
 *   * Neither the name of NVIDIA CORPORATION nor the names of its
 *     contributors may be used to endorse or promote products derived
 *     from this software without specific prior written permission.
 *
 * THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS ``AS IS'' AND ANY
 * EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE
 * IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR
 * PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT OWNER OR
 * CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL,
 * EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO,
 * PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR
 * PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY
 * OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
 * (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
 * OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
 */

/**
 * Vector addition: C = A + B.
```

```

*
* This sample is a very basic sample that implements element by element
* vector addition. It is the same as the sample illustrating Chapter 2
* of the programming guide with some additions like error checking.
*/

#include <stdio.h>

// For the CUDA runtime routines (prefixed with "cuda_")
#include <cuda_runtime.h>

#include <helper_cuda.h>
/**
 * CUDA Kernel Device code
 *
 * Computes the vector addition of A and B into C. The 3 vectors have the same
 * number of elements numElements.
 */
__global__ void vectorAdd(const float *A, const float *B, float *C,
                          int numElements) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;

    if (i < numElements) {
        C[i] = A[i] + B[i] + 0.0f;
    }
}

/**
 * Host main routine
 */
int main(void) {
    // Error code to check return values for CUDA calls
    cudaError_t err = cudaSuccess;

    for (int numElements = 10000000; numElements <= 100000000; numElements += 10000000){

        // CudaEvents for time measurements
        cudaEvent_t start, end;
        cudaEventCreate(&start);
        cudaEventCreate(&end);

        // Print the vector length to be used, and compute its size

        size_t size = numElements * sizeof(float);

        // Allocate the host input vector A

```

```

float *h_A = (float *)malloc(size);

// Allocate the host input vector B
float *h_B = (float *)malloc(size);

// Allocate the host output vector C
float *h_C = (float *)malloc(size);

// Verify that allocations succeeded
if (h_A == NULL || h_B == NULL || h_C == NULL) {
    fprintf(stderr, "Failed to allocate host vectors!\n");
    exit(EXIT_FAILURE);
}

// Initialize the host input vectors
for (int i = 0; i < numElements; ++i) {
    h_A[i] = rand() / (float)RAND_MAX;
    h_B[i] = rand() / (float)RAND_MAX;
}

// Allocate the device input vector A
float *d_A = NULL;
err = cudaMalloc((void **)&d_A, size);

if (err != cudaSuccess) {
    fprintf(stderr, "Failed to allocate device vector A (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

// Allocate the device input vector B
float *d_B = NULL;
err = cudaMalloc((void **)&d_B, size);

if (err != cudaSuccess) {
    fprintf(stderr, "Failed to allocate device vector B (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

// Allocate the device output vector C
float *d_C = NULL;
err = cudaMalloc((void **)&d_C, size);

if (err != cudaSuccess) {
    fprintf(stderr, "Failed to allocate device vector C (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

```

```

}

// Copy the host input vectors A and B in host memory to the device input
// vectors in
// device memory
err = cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);

if (err != cudaSuccess) {
    fprintf(stderr,
        "Failed to copy vector A from host to device (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

err = cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

if (err != cudaSuccess) {
    fprintf(stderr,
        "Failed to copy vector B from host to device (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

// Launch the Vector Add CUDA Kernel
for (int threadsPerBlock = 16; threadsPerBlock <= 256; threadsPerBlock *= 2){
    int blocksPerGrid = (numElements + threadsPerBlock - 1) / threadsPerBlock;

    for (int i = 0; i < 5; i++){
        cudaEventRecord(start);
        vectorAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, numElements);
        cudaEventRecord(end);
        cudaEventSynchronize(end);

        err = cudaGetLastError();

        if (err != cudaSuccess) {
            fprintf(stderr, "Failed to launch vectorAdd kernel (error code %s)!\n",
                cudaGetErrorString(err));
            exit(EXIT_FAILURE);
        }
    }

    // Copy the device result vector in device memory to the host result vector
    // in host memory.
    err = cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    if (err != cudaSuccess) {
        fprintf(stderr,
            "Failed to copy vector C from device to host (error code %s)!\n",

```

```

        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

// Verify that the result vector is correct
for (int i = 0; i < numElements; ++i) {
    if (fabs(h_A[i] + h_B[i] - h_C[i]) > 1e-5) {
        fprintf(stderr, "Result verification failed at element %d!\n", i);
        exit(EXIT_FAILURE);
    }
}

float time = 0;
cudaEventElapsedTime(&time, start, end);
printf("%d,%d,%f,%d\n", numElements, threadsPerBlock, time, i);
}
}

// Free device global memory
err = cudaFree(d_A);

if (err != cudaSuccess) {
    fprintf(stderr, "Failed to free device vector A (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

err = cudaFree(d_B);

if (err != cudaSuccess) {
    fprintf(stderr, "Failed to free device vector B (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

err = cudaFree(d_C);

if (err != cudaSuccess) {
    fprintf(stderr, "Failed to free device vector C (error code %s)!\n",
        cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}

// Free host memory
free(h_A);
free(h_B);
free(h_C);
}

printf("Done\n");

```

```

    return 0;
}

```

```

[53]: import pandas as pd
import seaborn as sns

sns.set(rc={'figure.figsize':(11.7,8.27)})

```

```

[48]: data = pd.read_csv("data/1.csv")
data.head()

```

```

[48]:
  number_of_elements  threads_per_block    time  iter
0          100000000             16  1.949024     0
1          100000000             16  1.958880     1
2          100000000             16  1.961888     2
3          100000000             16  1.954432     3
4          100000000             16  1.955776     4

```

```


[56]: sns.lineplot(data=data, x="number_of_elements", y="time",
    ↪hue="threads_per_block", marker="o").set(xlabel="Number of elements",
    ↪ylabel="Time [ms]", title="Time per number of elements")

```

```

[56]: [Text(0.5, 0, 'Number of elements'),
Text(0, 0.5, 'Time [ms]'),
Text(0.5, 1.0, 'Time per number of elements')]

```



1.2 Ćw 2

Kod naiwny:

```
%%cu
#include<stdio.h>
#include<stdlib.h>

#define BLOCK_SIZE 32

__global__ void matrix_transpose_naive(int *input, int *output, int n) {

    int indexX = threadIdx.x + blockIdx.x * blockDim.x;
    int indexY = threadIdx.y + blockIdx.y * blockDim.y;
    int index = indexY * n + indexX;
    int transposedIndex = indexX * n + indexY;

    // this has discoalesced global memory store
    output[transposedIndex] = input[index];

    // this has discoalesced global memore load
    // output[index] = input[transposedIndex];
}

__global__ void matrix_transpose_shared(int *input, int *output, int n) {

    __shared__ int sharedMemory [BLOCK_SIZE] [BLOCK_SIZE];

    // global index
    int indexX = threadIdx.x + blockIdx.x * blockDim.x;
    int indexY = threadIdx.y + blockIdx.y * blockDim.y;

    // transposed global memory index
    int tindexX = threadIdx.x + blockIdx.y * blockDim.x;
    int tindexY = threadIdx.y + blockIdx.x * blockDim.y;

    // local index
    int localIndexX = threadIdx.x;
    int localIndexY = threadIdx.y;

    int index = indexY * n + indexX;
    int transposedIndex = tindexY * n + tindexX;

    // reading from global memory in coalesed manner and performing tanspose in shared memory
    sharedMemory[localIndexX][localIndexY] = input[index];
```

```

    __syncthreads();

    // writing into global memory in coalesced fashion via transposed data in shared memory
    output[transposedIndex] = sharedMemory[localIndexY][localIndexX];
}

//basically just fills the array with index.
void fill_array(int *data, int n) {
    for(int idx=0;idx<(n*n);idx++)
        data[idx] = idx;
}

void print_output(int *a, int *b, int n) {
    printf("\n Original Matrix::\n");
    for(int idx=0;idx<(n*n);idx++) {
        if(idx%n == 0)
            printf("\n");
        printf(" %d ", a[idx]);
    }
    printf("\n Transposed Matrix::\n");
    for(int idx=0;idx<(n*n);idx++) {
        if(idx%n == 0)
            printf("\n");
        printf(" %d ", b[idx]);
    }
}

int main(void) {
    int *a, *b;
    int *d_a, *d_b; // device copies of a, b, c

    cudaEvent_t start, end;
    cudaEventCreate(&start);
    cudaEventCreate(&end);

    for (int n = 1028; n <= 10280; n += 1028){
        int size = n * n * sizeof(int);

        // Alloc space for host copies of a, b, c and setup input values
        a = (int *)malloc(size); fill_array(a, n);
        b = (int *)malloc(size);

        // Alloc space for device copies of a, b, c
        cudaMalloc((void **)&d_a, size);
        cudaMalloc((void **)&d_b, size);

        // Copy inputs to device

```



```

    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

    dim3 blockSize(BLOCK_SIZE,BLOCK_SIZE,1);
    dim3 gridSize(n/BLOCK_SIZE,n/BLOCK_SIZE,1);

    for (int i = 0; i < 5; i++){
        cudaEventRecord(start);
        // matrix_transpose_naive<<<gridSize,blockSize>>>(d_a,d_b,n);
        matrix_transpose_shared<<<gridSize,blockSize>>>(d_a,d_b,n);
        cudaEventRecord(end);
        cudaEventSynchronize(end);

        float time = 0;
        cudaEventElapsedTime(&time, start, end);
        printf("%d,%f,%d\n", n, time, i);
    }

    free(a);
    free(b);
    cudaFree(d_a);
    cudaFree(d_b);
}

return 0;
}

```

Kod wykorzystujący pamięć współdzieloną:

```

%%cu
#include<stdio.h>
#include<stdlib.h>

#define BLOCK_SIZE 32

__global__ void matrix_transpose_naive(int *input, int *output, int n) {

    int indexX = threadIdx.x + blockIdx.x * blockDim.x;
    int indexY = threadIdx.y + blockIdx.y * blockDim.y;
    int index = indexY * n + indexX;
    int transposedIndex = indexX * n + indexY;

    // this has discoalesced global memory store
    output[transposedIndex] = input[index];

    // this has discoalesced global memore load
    // output[index] = input[transposedIndex];
}

__global__ void matrix_transpose_shared(int *input, int *output, int n) {

```

```

__shared__ int sharedMemory [BLOCK_SIZE] [BLOCK_SIZE];

// global index
int indexX = threadIdx.x + blockIdx.x * blockDim.x;
int indexY = threadIdx.y + blockIdx.y * blockDim.y;

// transposed global memory index
int tindexX = threadIdx.x + blockIdx.y * blockDim.x;
int tindexY = threadIdx.y + blockIdx.x * blockDim.y;

// local index
int localIndexX = threadIdx.x;
int localIndexY = threadIdx.y;

int index = indexY * n + indexX;
int transposedIndex = tindexY * n + tindexX;

// reading from global memory in coalesced manner and performing transpose in shared memory
sharedMemory[localIndexX][localIndexY] = input[index];

__syncthreads();

// writing into global memory in coalesced fashion via transposed data in shared memory
output[transposedIndex] = sharedMemory[localIndexY][localIndexX];
}

//basically just fills the array with index.
void fill_array(int *data, int n) {
    for(int idx=0;idx<(n*n);idx++)
        data[idx] = idx;
}

void print_output(int *a, int *b, int n) {
    printf("\n Original Matrix::\n");
    for(int idx=0;idx<(n*n);idx++) {
        if(idx%n == 0)
            printf("\n");
        printf(" %d ", a[idx]);
    }
    printf("\n Transposed Matrix::\n");
    for(int idx=0;idx<(n*n);idx++) {
        if(idx%n == 0)
            printf("\n");
        printf(" %d ", b[idx]);
    }
}

```

```

int main(void) {
    int *a, *b;
    int *d_a, *d_b; // device copies of a, b, c

    cudaEvent_t start, end;
    cudaEventCreate(&start);
    cudaEventCreate(&end);

    for (int n = 1028; n <= 10280; n += 1028){
        int size = n * n * sizeof(int);

        // Alloc space for host copies of a, b, c and setup input values
        a = (int *)malloc(size); fill_array(a, n);
        b = (int *)malloc(size);

        // Alloc space for device copies of a, b, c
        cudaMalloc((void **)&d_a, size);
        cudaMalloc((void **)&d_b, size);

        // Copy inputs to device
        cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
        cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

        dim3 blockSize(BLOCK_SIZE, BLOCK_SIZE, 1);
        dim3 gridSize(n/BLOCK_SIZE, n/BLOCK_SIZE, 1);

        for (int i = 0; i < 5; i++){
            cudaEventRecord(start);
            matrix_transpose_naive<<<gridSize, blockSize>>>(d_a, d_b, n);
            // matrix_transpose_shared<<<gridSize, blockSize>>>(d_a, d_b, n);
            cudaEventRecord(end);
            cudaEventSynchronize(end);

            float time = 0;
            cudaEventElapsedTime(&time, start, end);
            printf("%d,%f,%d\n", n, time, i);
        }

        free(a);
        free(b);
        cudaFree(d_a);
        cudaFree(d_b);
    }

    return 0;
}


```

```
[70]: data2 = pd.read_csv("data/2.csv")
```

```
[71]: ax = sns.lineplot(data=data2, y="t", x="n", hue="l", marker="o")
ax.set(title="Naive vs Shared", xlabel="Matrix width and lenght", ylabel="Time_
↪ [ms]")
ax.legend(title="Type of memory")

ax
```

```
[71]: <AxesSubplot:title={'center':'Naive vs Shared'}, xlabel='Matrix width and
length', ylabel='Time [ms] '>
```



```
[ ]:
```